

## Basic info about running platform

Instead of copy pasting the same machine specification we'll just specify the name of the machine next to the results

### System - "Main machine"

- Time of this report: 9/18/2018, 17:43:26
- Operating System: Windows 10 Pro 64-bit
- System Manufacturer: Dell Inc.
- System Model: Inspiron 7559
- BIOS: 1.2.2 (type: UEFI)
- Processor: Intel(R) Core(TM) i7-6700HQ CPU @ 2.60GHz (8 CPUs), ~2.6GHz
- Memory: 16384MB RAM
- Available OS Memory: 16256MB RAM

### System - "Jet Engine"

- Date: 2018-09-17T17:39:20+0200
- OS: Windows 10; 10.0; amd64
- JVM: Oracle Corporation; 10.0.2
- CPU: Intel64 Family 6 Model 60 Stepping 3, GenuineIntel; 8 "cores"
- RAM: 16 GB

## Exercise 3.1

### Part 1.

#### Mark01

As written in the in the article the parallelism performed by the CPU lowered the result time into `0.0 ns` . We have decided to lower the number of executions to see whether we get similar results => and indeed we do get similarities described in the benchmarking

```
-----  
|Run on - "Main machine"|  
-----  
100 => `119.0 ns`  
100_000 => `29.9 ns`  
100_000_000 runs => `0.0 ns`
```

#### Mark02

We can see that our result time is slightly lower then the one in `benchmarking` , which is cause (we assume) by having a better platform.

```
-----  
|Run on - "Main machine"|  
-----  
results => `23.2 ns`
```

Mark03

We can observe the same consistency and similarity with Mark02 , however the CPU time returned seems to stabilize around 25.5 ns

```
-----  
|Run on - "Main machine"|  
-----  
23.2 ns  
25.6 ns  
25.6 ns  
26.4 ns  
26.1 ns  
25.5 ns  
25.4 ns  
25.4 ns  
25.5 ns  
25.5 ns
```

Mark04

The result time is slightly lower than the one in Mark02 , however the consistency still lingers between 21 - 25 ns

```
-----  
|Run on - "Main machine"|  
-----  
results => `21.2 ns +/- 0.503`
```

Mark05

We can observe similarly fashioned results as in the benchmarking . We can conclude that test from Mark01 seems to be proven in a nice list.

```
-----  
|Run on - "Main machine"|  
-----  
345.0 ns +/- 600.21      2  
137.5 ns +/- 84.37      4  
147.5 ns +/- 149.98     8  
137.5 ns +/- 154.48    16  
35.6 ns +/- 7.40       32  
65.0 ns +/- 70.32      64  
76.6 ns +/- 59.48     128  
31.7 ns +/- 2.04       256  
39.8 ns +/- 11.93      512  
40.5 ns +/- 8.03       1024  
32.7 ns +/- 3.52       2048  
32.9 ns +/- 4.06       4096  
32.9 ns +/- 3.34       8192  
33.2 ns +/- 4.74      16384  
31.2 ns +/- 17.39     32768  
29.3 ns +/- 5.57     65536  
21.7 ns +/- 4.72    131072  
20.7 ns +/- 2.11   262144  
20.2 ns +/- 1.39   524288  
20.3 ns +/- 0.79  1048576  
20.2 ns +/- 0.66 2097152
```

```

20.1 ns +/-      0.25      4194304
19.9 ns +/-      0.25      8388608
19.9 ns +/-      0.07      16777216

```

## Mark06

We can observe that our running platform was able to perform up to 16,7 million runs with average 22.6 ns per call and the deviation between 10 calls is 0.14

```

-----
|Run on - "Main machine"|
-----
multiply          510.0 ns      1069.22      2
multiply          575.0 ns      1194.78      4
multiply          166.3 ns      112.28       8
multiply          271.3 ns      348.24      16
multiply           72.2 ns       94.34      32
multiply          49.7 ns       21.79      64
multiply          91.3 ns       69.49     128
multiply          43.6 ns        9.65     256
multiply          65.4 ns       11.98     512
multiply          37.1 ns        4.56    1024
multiply          36.0 ns        2.27    2048
multiply          38.7 ns        3.62    4096
multiply          30.2 ns        3.01    8192
multiply          30.4 ns        1.84   16384
multiply          27.1 ns        4.01   32768
multiply          37.9 ns       14.77   65536
multiply          24.0 ns        3.28  131072
multiply          24.0 ns        1.26  262144
multiply          23.2 ns        1.68  524288
multiply          23.0 ns        1.21 1048576
multiply          22.9 ns        0.51 2097152
multiply          22.6 ns        0.37 4194304
multiply          22.7 ns        0.31 8388608
multiply          22.6 ns        0.14 16777216

```

## Part 2.

Considering part one of this exercise the results seem plausible.

- `pow` => our execution platform made 4 times more calls with lower time results between the calls but higher deviation
- `acos` => result time between the calls is similar as one in the `benchmarking` however deviation is waay higher. This may be affected by number of calls made, which for our platform is 2 times lower.
- `atan` => same number of calls, somewhat close result time but again deviation is waaay higher

```

-----
|Run on - "Main machine"|
-----
pow              23.4 ns       1.27   16777216
exp              23.6 ns       0.15   16777216
log              12.4 ns       0.06   33554432
sin              16.4 ns       0.40   16777216
cos              15.8 ns       0.52   16777216
tan              21.9 ns       0.69   16777216
asin            232.8 ns       3.95   2097152
acos            230.2 ns      15.73  1048576
atan            47.2 ns       1.47   8388608

```

## Exercise 3.2

## Part 1.

- hashCode() => result time got very fast and deviation very small with increased number of calls
- Point creation => by group's experience, we expected this to be heavier operation, which is proven by relatively high result time 64.1 ns and deviation 1.15 for 4194304 calls
- Thread's work => this is expected to have least runtime from all of the "threads" measurements except creation of the thread.
- Thread create => relatively heavy operation compared to let's say point creation
- Thread create start => we did not expected that the starting of the thread would increase result time and deviation so drastically.
- Thread create start join => this was expected to take the most work to complete as it performs all of the "thread" measurements before
- Uncontended lock => relatively light operation, which seems to be really consistent after couple of calls all the way until the end of the measurement

## Part 2.

```

-----
|Run on - "Main machine"|
-----
hashCode()                2.7 ns      0.02  134217728
Point creation             65.4 ns      0.87   4194304
Thread's work             5400.8 ns     45.99   65536
Thread create             911.3 ns     111.43  524288
Thread create start       95256.7 ns   3952.53   4096
Thread create start join  194481.8 ns  7644.20   2048
Uncontended lock          5.5 ns      0.52  67108864

```

```

-----
|Run on - "Jet Engine"|
-----
hashCode()                2.4 ns      0.00  134217728
Point creation             62.6 ns      0.26   4194304
Thread's work             5042.3 ns     4.02   65536
Thread create             818.1 ns      99.30  524288
Thread create start      137040.2 ns   593.13   2048
Thread create start join  232831.9 ns  599.28   2048
ai value = 1392580000
Uncontended lock          4.4 ns      0.61  67108864

```

## Exercise 3.3

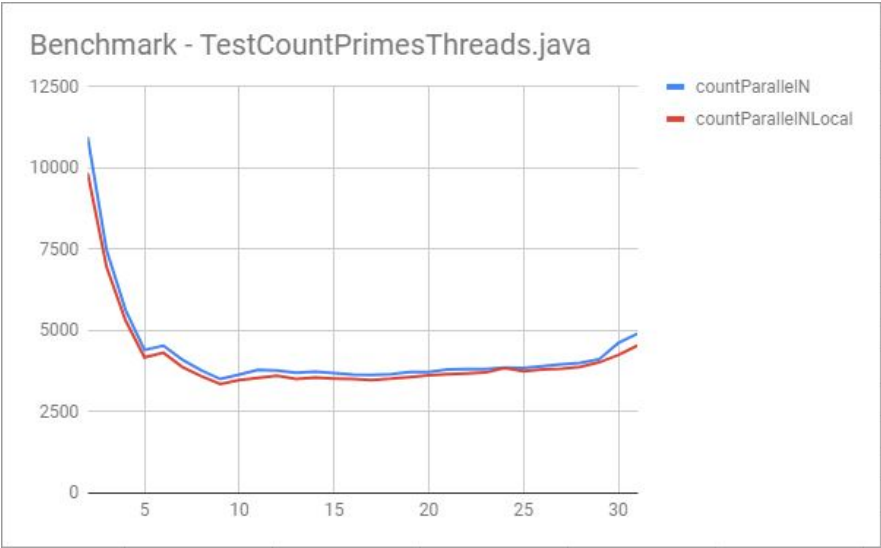
### Part 1.

countSequential		9717.4 us	324.81	32
countParallelN	1	10959.7 us	255.06	32
countParallelNLocal	1	9847.3 us	150.60	32
countParallelN	2	7470.9 us	114.03	64
countParallelNLocal	2	6952.6 us	161.13	64
countParallelN	3	5634.6 us	72.63	64
countParallelNLocal	3	5312.1 us	79.59	64
countParallelN	4	4403.8 us	117.83	64
countParallelNLocal	4	4173.0 us	144.46	64
countParallelN	5	4533.0 us	107.24	64
countParallelNLocal	5	4315.7 us	113.51	64
countParallelN	6	4103.5 us	80.75	64
countParallelNLocal	6	3875.4 us	84.33	128

countParallelN	7	3771.2 us	136.94	128
countParallelNLocal	7	3594.8 us	52.97	128
countParallelN	8	3507.9 us	51.82	128
countParallelNLocal	8	3351.9 us	40.23	128
countParallelN	9	3646.7 us	65.22	128
countParallelNLocal	9	3478.8 us	43.97	128
countParallelN	10	3794.4 us	137.67	128
countParallelNLocal	10	3546.3 us	52.40	128
countParallelN	11	3768.3 us	48.85	128
countParallelNLocal	11	3611.8 us	59.16	128
countParallelN	12	3705.7 us	39.28	128
countParallelNLocal	12	3514.2 us	53.43	128
countParallelN	13	3740.7 us	41.33	128
countParallelNLocal	13	3554.6 us	42.14	128
countParallelN	14	3690.9 us	49.24	128
countParallelNLocal	14	3527.6 us	38.33	128
countParallelN	15	3646.1 us	42.41	128
countParallelNLocal	15	3511.2 us	49.41	128
countParallelN	16	3639.1 us	70.27	128
countParallelNLocal	16	3480.2 us	52.71	128
countParallelN	17	3653.9 us	55.93	128
countParallelNLocal	17	3517.6 us	71.03	128
countParallelN	18	3719.4 us	53.60	128
countParallelNLocal	18	3570.9 us	47.53	128
countParallelN	19	3719.9 us	35.17	128
countParallelNLocal	19	3618.5 us	62.79	128
countParallelN	20	3806.9 us	114.98	64
countParallelNLocal	20	3652.8 us	83.87	128
countParallelN	21	3809.3 us	51.35	128
countParallelNLocal	21	3681.2 us	69.94	128
countParallelN	22	3820.1 us	57.93	128
countParallelNLocal	22	3708.4 us	51.48	128
countParallelN	23	3856.7 us	78.59	128
countParallelNLocal	23	3845.4 us	250.96	128
countParallelN	24	3848.8 us	63.68	128
countParallelNLocal	24	3748.2 us	91.82	64
countParallelN	25	3901.8 us	61.90	128
countParallelNLocal	25	3807.0 us	65.78	128
countParallelN	26	3961.5 us	80.46	64
countParallelNLocal	26	3825.6 us	89.45	128
countParallelN	27	4002.4 us	88.58	64
countParallelNLocal	27	3880.4 us	71.93	128
countParallelN	28	4123.5 us	181.32	64
countParallelNLocal	28	4025.4 us	124.40	64
countParallelN	29	4618.1 us	401.48	64
countParallelNLocal	29	4248.9 us	298.49	64
countParallelN	30	4903.0 us	311.60	64
countParallelNLocal	30	4536.0 us	188.44	64
countParallelN	31	4506.1 us	172.48	64
countParallelNLocal	31	4773.0 us	295.86	64
countParallelN	32	4607.6 us	164.94	64
countParallelNLocal	32	4472.4 us	143.61	64

## Part 2.

---

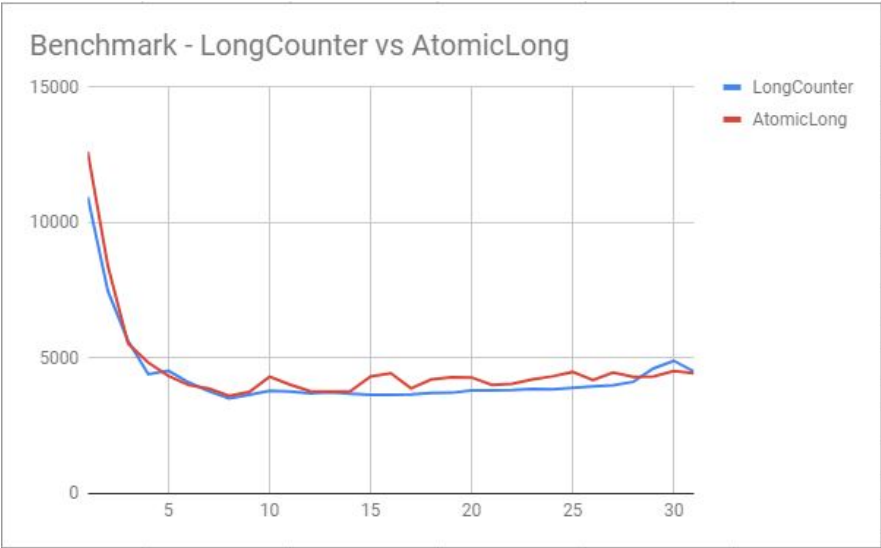


Part 3.

The results seem plausible. We can see that lowest runtime was with 8 threads which nicely corresponds to the number of cores on the platform.

Part 4.

It does not seem that the AtomicLong performed much better. One should generally use build in helpers as they are usually most optimized.




countSequential		11343.8 us	210.28	32
countParallelN	1	12621.9 us	421.09	32
countParallelNLocal	1	12091.0 us	655.20	32
countParallelN	2	8397.6 us	501.94	32
countParallelNLocal	2	7687.8 us	326.20	64
countParallelN	3	5537.6 us	136.25	64
countParallelNLocal	3	5764.6 us	104.79	64
countParallelN	4	4834.7 us	157.98	64
countParallelNLocal	4	5107.7 us	86.98	64
countParallelN	5	4343.1 us	74.64	64
countParallelNLocal	5	4521.6 us	80.69	64
countParallelN	6	4002.9 us	76.58	64
countParallelNLocal	6	4317.2 us	150.01	64
countParallelN	7	3872.9 us	133.35	64
countParallelNLocal	7	4018.6 us	164.85	64

countParallelN	8	3610.9 us	70.10	128
countParallelNLocal	8	3680.1 us	67.85	128
countParallelN	9	3753.9 us	75.18	128
countParallelNLocal	9	3844.6 us	59.19	128
countParallelN	10	4312.6 us	413.00	64
countParallelNLocal	10	4137.2 us	208.71	64
countParallelN	11	4025.7 us	245.90	64
countParallelNLocal	11	4116.4 us	300.70	64
countParallelN	12	3783.7 us	30.03	128
countParallelNLocal	12	3870.5 us	40.87	128
countParallelN	13	3766.1 us	39.25	128
countParallelNLocal	13	3839.1 us	45.61	64
countParallelN	14	3774.5 us	89.95	128
countParallelNLocal	14	3757.2 us	43.88	128
countParallelN	15	4322.1 us	460.74	64
countParallelNLocal	15	4729.8 us	516.31	64
countParallelN	16	4445.2 us	425.57	128
countParallelNLocal	16	4046.9 us	206.38	64
countParallelN	17	3883.7 us	87.16	64
countParallelNLocal	17	4162.2 us	230.79	64
countParallelN	18	4212.0 us	343.34	128
countParallelNLocal	18	4070.3 us	100.67	64
countParallelN	19	4290.9 us	188.31	64
countParallelNLocal	19	4153.4 us	289.48	64
countParallelN	20	4284.3 us	204.46	64
countParallelNLocal	20	4292.5 us	224.66	64
countParallelN	21	4013.2 us	121.64	64
countParallelNLocal	21	4070.3 us	89.31	64
countParallelN	22	4055.9 us	112.56	64
countParallelNLocal	22	4172.0 us	210.07	64
countParallelN	23	4217.8 us	242.38	64
countParallelNLocal	23	4364.8 us	222.57	64
countParallelN	24	4328.0 us	223.61	64
countParallelNLocal	24	4403.0 us	230.42	64
countParallelN	25	4488.0 us	334.32	64
countParallelNLocal	25	4232.7 us	134.25	64
countParallelN	26	4194.4 us	142.11	64
countParallelNLocal	26	4387.9 us	145.40	64
countParallelN	27	4462.9 us	202.43	64
countParallelNLocal	27	4650.8 us	332.43	64
countParallelN	28	4310.0 us	100.38	64
countParallelNLocal	28	4361.1 us	205.17	64
countParallelN	29	4320.8 us	123.94	64
countParallelNLocal	29	4847.2 us	1172.27	32
countParallelN	30	4519.4 us	163.24	64
countParallelNLocal	30	4665.0 us	552.00	64
countParallelN	31	4446.5 us	100.61	64
countParallelNLocal	31	4580.3 us	171.59	64
countParallelN	32	4536.3 us	95.21	64
countParallelNLocal	32	4566.1 us	135.05	64

## Part 5.

---

Without count and With count



It seems to be slower.

countSequential		11617.1 us	280.84	32
countParallelN	1	12956.4 us	358.59	32
countParallelNLocal	1	12362.5 us	615.54	32
countParallelN	2	8203.4 us	131.86	32
countParallelNLocal	2	7911.9 us	302.01	64
countParallelN	3	6117.9 us	323.20	64
countParallelNLocal	3	5968.3 us	342.40	64
countParallelN	4	5136.8 us	125.38	64
countParallelNLocal	4	5009.9 us	81.44	64
countParallelN	5	4620.9 us	149.78	64
countParallelNLocal	5	4620.1 us	138.86	64
countParallelN	6	4188.6 us	56.08	64
countParallelNLocal	6	4170.0 us	62.93	64
countParallelN	7	4124.4 us	135.99	64
countParallelNLocal	7	3928.8 us	103.28	64
countParallelN	8	3804.4 us	88.44	128
countParallelNLocal	8	4433.5 us	745.28	64
countParallelN	9	4913.2 us	240.33	64
countParallelNLocal	9	4800.9 us	464.80	64
countParallelN	10	4115.1 us	229.96	64
countParallelNLocal	10	4172.8 us	231.16	64
countParallelN	11	4076.7 us	89.61	64
countParallelNLocal	11	4183.2 us	158.18	64
countParallelN	12	4147.2 us	200.84	128
countParallelNLocal	12	4241.6 us	233.07	64
countParallelN	13	4129.7 us	232.57	64
countParallelNLocal	13	4153.0 us	257.66	64
countParallelN	14	4247.3 us	327.65	128
countParallelNLocal	14	4226.6 us	271.72	64
countParallelN	15	3926.1 us	75.72	64
countParallelNLocal	15	4156.2 us	257.04	64
countParallelN	16	4218.3 us	244.50	64
countParallelNLocal	16	3901.1 us	113.59	64
countParallelN	17	4219.9 us	258.01	64
countParallelNLocal	17	4113.3 us	259.23	64
countParallelN	18	4936.5 us	682.63	64
countParallelNLocal	18	4638.1 us	745.18	64
countParallelN	19	4550.1 us	167.57	64
countParallelNLocal	19	4436.6 us	200.99	64
countParallelN	20	4738.8 us	689.16	64
countParallelNLocal	20	4455.9 us	336.87	64
countParallelN	21	4119.1 us	91.59	64
countParallelNLocal	21	4182.1 us	137.57	64
countParallelN	22	4190.0 us	139.34	64
countParallelNLocal	22	4266.5 us	184.64	64
countParallelN	23	4230.9 us	82.57	64
countParallelNLocal	23	4264.9 us	168.74	64
countParallelN	24	4304.9 us	200.03	64



countParallelNLocal	24	4258.5 us	204.91	64
countParallelN	25	4368.0 us	220.70	64
countParallelNLocal	25	4494.0 us	382.03	64
countParallelN	26	4528.5 us	311.13	64
countParallelNLocal	26	4448.4 us	169.13	64
countParallelN	27	4806.1 us	402.12	64
countParallelNLocal	27	4628.5 us	266.74	64
countParallelN	28	4438.3 us	214.75	64
countParallelNLocal	28	4637.2 us	168.85	64
countParallelN	29	5711.2 us	638.89	64
countParallelNLocal	29	4602.2 us	123.57	64
countParallelN	30	4638.7 us	69.37	64
countParallelNLocal	30	4915.0 us	716.66	64
countParallelN	31	4736.1 us	314.77	64
countParallelNLocal	31	5021.4 us	505.33	64
countParallelN	32	4692.3 us	114.92	64
countParallelNLocal	32	4763.6 us	129.21	64

## Exercise 3.4

---

### Part 1.

---

Memoizer1 => Results 3.9 s 0.12 2

### Part 2.

---

Memoizer2 => Results 2.0 s 0.05 2

### Part 3.

---

Memoizer3 => Results 1.8 s 0.04 2

### Part 4.

---

Memoizer4 => Results 1.8 s 0.04 2

### Part 5.

---

Memoizer5 => Results 1.7 s 0.03 2

### Part 6.

---

Memoizer0 => Results 1.8 s 0.03 2

### Part 7.

---

Memoizer5 performed the best, with checking the cache if it exists before saving it. The results agree with lecture's and Goetz's development of the cache classes.

### Part 8.

---

The current tests are running on empty caches, so we were thinking on running the tests for 3 different stages of cache - empty, randomized, predefined. This way we can have more data on how are the cache implementations running with either empty or filled storages.

